

## Project Code: For Different Tree Boosts

### Comparative analysis of different tree Boosts for IDS on Intrusion Detection

#### Tree Boosts to Be Compared:

##### 1. XGB Boost

##### 2. LGB Boost

##### 3. Cat Boost

##### 4. Extra Trees

##### 5. Random forests

##### 6. Decision trees

#### Import Libraries

1. Pandas
2. Numpy
3. matplotlib
4. Libraries for different Tree Models ( LightGBM,xgBoost,CatBoost)
5. other libraries as needed.

```
import warnings
warnings.filterwarnings("ignore")

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.metrics import
classification_report,confusion_matrix,accuracy_score,
precision_score, recall_score, f1_score
import lightgbm as lgb
!pip3 install catboost
```

```
import catboost as cbt
import xgboost as xgb
import time
!pip3 install river
from river import stream
from statistics import mode
```

Looking in indexes: <https://pypi.org/simple>, <https://us-python.pkg.dev/colab-wheels/public/simple/>

Collecting catboost

Downloading catboost-1.1.1-cp38-none-manylinux1\_x86\_64.whl (76.6 MB)  
Requirement already satisfied: scipy in /usr/local/lib/python3.8/dist-packages (from catboost) (1.7.3)

Requirement already satisfied: matplotlib in /usr/local/lib/python3.8/dist-packages (from catboost) (3.2.2)

Requirement already satisfied: plotly in /usr/local/lib/python3.8/dist-packages (from catboost) (5.5.0)

Requirement already satisfied: graphviz in /usr/local/lib/python3.8/dist-packages (from catboost) (0.10.1)

Requirement already satisfied: pandas>=0.24.0 in /usr/local/lib/python3.8/dist-packages (from catboost) (1.3.5)

Requirement already satisfied: six in /usr/local/lib/python3.8/dist-packages (from catboost) (1.15.0)

Requirement already satisfied: numpy>=1.16.0 in /usr/local/lib/python3.8/dist-packages (from catboost) (1.21.6)

Requirement already satisfied: python-dateutil>=2.7.3 in /usr/local/lib/python3.8/dist-packages (from pandas>=0.24.0->catboost) (2.8.2)

Requirement already satisfied: pytz>=2017.3 in /usr/local/lib/python3.8/dist-packages (from pandas>=0.24.0->catboost) (2022.6)

Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.8/dist-packages (from matplotlib->catboost) (0.11.0)

Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.8/dist-packages (from matplotlib->catboost) (1.4.4)

Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /usr/local/lib/python3.8/dist-packages (from matplotlib->catboost) (3.0.9)

Requirement already satisfied: tenacity>=6.2.0 in /usr/local/lib/python3.8/dist-packages (from plotly->catboost) (8.1.0)

Installing collected packages: catboost

Successfully installed catboost-1.1.1

Looking in indexes: <https://pypi.org/simple>, <https://us-python.pkg.dev/colab-wheels/public/simple/>

Collecting river

Downloading river-0.14.0-cp38-cp38-manylinux\_2\_17\_x86\_64.manylinux2014\_x86\_64.whl (3.2 MB)

Requirement already satisfied: scipy>=1.5 in /usr/local/lib/python3.8/dist-

```

packages (from river) (1.7.3)
Requirement already satisfied: pandas>=1.3 in
/usr/local/lib/python3.8/dist-packages (from river) (1.3.5)
Collecting numpy>=1.23.4
  Downloading numpy-1.24.1-cp38-cp38-
manylinux_2_17_x86_64.manylinux2014_x86_64.whl (17.3 MB)
Requirement already satisfied: pytz>=2017.3 in /usr/local/lib/python3.8/dist-
packages (from pandas>=1.3->river) (2022.6)
Requirement already satisfied: python-dateutil>=2.7.3 in
/usr/local/lib/python3.8/dist-packages (from pandas>=1.3->river)
(2.8.2)
Requirement already satisfied: six>=1.5 in
/usr/local/lib/python3.8/dist-packages (from python-dateutil>=2.7.3-
>pandas>=1.3->river) (1.15.0)
Collecting scipy>=1.5
  Downloading scipy-1.9.3-cp38-cp38-
manylinux_2_17_x86_64.manylinux2014_x86_64.whl (33.8 MB)
py, scipy, river
  Attempting uninstall: numpy
    Found existing installation: numpy 1.21.6
    Uninstalling numpy-1.21.6:
      Successfully uninstalled numpy-1.21.6
  Attempting uninstall: scipy
    Found existing installation: scipy 1.7.3
    Uninstalling scipy-1.7.3:
      Successfully uninstalled scipy-1.7.3
ERROR: pip's dependency resolver does not currently take into account
all the packages that are installed. This behaviour is the source of
the following dependency conflicts.
numba 0.56.4 requires numpy<1.24,>=1.18, but you have numpy 1.24.1
which is incompatible.
Successfully installed numpy-1.24.1 river-0.14.0 scipy-1.9.3

```

```

{"pip_warning":{"packages":["numpy"]}}

```

```

from google.colab import drive
drive.mount('/content/drive')

```

Mounted at /content/drive

## DATA SET TO BE USED: CICIDS2017 dataset

```

df = pd.read_csv("/content/drive/MyDrive/ISM/Intrusion-Detection-
System-Using-Machine-Learning-main/data/CICIDS2017_sample_km.csv")

```

```

df.Label.value_counts()

```

```

0    18225
3     3042
6     2180
1     1966

```

```

5      1255
2        96
4        36
Name: Label, dtype: int64

```

```
df.head()
```

	Flow Duration	Total Fwd Packets	Total Backward Packets	\
0	5.416666e-07	0.0	0.000003	
1	5.416666e-07	0.0	0.000003	
2	4.416666e-07	0.0	0.000003	
3	7.499999e-07	0.0	0.000003	
4	7.249999e-07	0.0	0.000003	

	Total Length of Fwd Packets	Total Length of Bwd Packets	\
0	4.651163e-07	9.153974e-09	
1	4.651163e-07	9.153974e-09	
2	4.651163e-07	9.153974e-09	
3	4.651163e-07	9.153974e-09	
4	4.651163e-07	9.153974e-09	

	Fwd Packet Length Max	Fwd Packet Length Min	Fwd Packet Length
Mean \			
0	0.000242	0.002581	
0.00101			
1	0.000242	0.002581	
0.00101			
2	0.000242	0.002581	
0.00101			
3	0.000242	0.002581	
0.00101			
4	0.000242	0.002581	
0.00101			

	Fwd Packet Length Std	Bwd Packet Length Max	...
min_seg_size_forward \			
0	0.0	0.000307	...
1.0			
1	0.0	0.000307	...
1.0			
2	0.0	0.000307	...
1.0			
3	0.0	0.000307	...
1.0			
4	0.0	0.000307	...
1.0			

	Active Mean	Active Std	Active Max	Active Min	Idle Mean	Idle
Std \						
0	0.0	0.0	0.0	0.0	0.0	

```

0.0
1      0.0      0.0      0.0      0.0      0.0
0.0
2      0.0      0.0      0.0      0.0      0.0
0.0
3      0.0      0.0      0.0      0.0      0.0
0.0
4      0.0      0.0      0.0      0.0      0.0
0.0

```

```

      Idle Max  Idle Min  Label
0      0.0      0.0      0
1      0.0      0.0      0
2      0.0      0.0      0
3      0.0      0.0      0
4      0.0      0.0      0

```

[5 rows x 78 columns]

## Corresponding Attack Types:

1. 0 BENIGN 18225
2. 1 Bot 1966
3. 2 BruteForce 96
4. 3 DoS 3042
5. 4 Infiltration 36
6. 5 PortScan 1255
7. 6 WebAttack 2180

## Splitting train set and test set for training purpose.

```
train_size = 0.8, test_size = 0.2
```

```
X = df.drop(['Label'],axis=1)
```

```
y = df['Label']
```

```
X_train, X_test, y_train, y_test = train_test_split(X,y, train_size =
0.8, test_size = 0.2, random_state = 0) #shuffle=False
```

## SMOTE to solve class-imbalance

```
pd.Series(y_train).value_counts()
```

```

0      14569
3       2430
6       1728
1       1579
5       1024

```

```

2      82
4      28
Name: Label, dtype: int64

from imblearn.over_sampling import SMOTE
smote=SMOTE(n_jobs=-1,sampling_strategy={2:1000,4:1000})

X_train, y_train = smote.fit_resample(X_train, y_train)

pd.Series(y_train).value_counts()

0      14569
3       2430
6       1728
1       1579
5       1024
2       1000
4       1000
Name: Label, dtype: int64

```

## Machine Learning (ML) model training

Training three base learners: LightGBM, XGBoost, CatBoost

```


%%time
# Train the LightGBM algorithm
import lightgbm as lgb
lg = lgb.LGBMClassifier()
lg.fit(X_train, y_train)
y_pred = lg.predict(X_test)
print(classification_report(y_test,y_pred))
print("Accuracy of LightGBM: "+ str(accuracy_score(y_test, y_pred)))
print("Precision of LightGBM: "+ str(precision_score(y_test, y_pred,
average='weighted'))))
print("Recall of LightGBM: "+ str(recall_score(y_test, y_pred,
average='weighted'))))
print("Average F1 of LightGBM: "+ str(f1_score(y_test, y_pred,
average='weighted'))))
print("F1 of LightGBM for each type of attack: "+ str(f1_score(y_test,
y_pred, average=None)))
lg_f1=f1_score(y_test, y_pred, average=None)

# Plot the confusion matrix
cm=confusion_matrix(y_test,y_pred)
f,ax=plt.subplots(figsize=(5,5))
sns.heatmap(cm,annot=True,linewidth=0.5,linecolor="red",fmt=".0f",ax=ax)
plt.xlabel("y_pred")
plt.ylabel("y_true")
plt.show()

```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	3656
1	1.00	0.99	0.99	387
2	1.00	1.00	1.00	14
3	1.00	1.00	1.00	612
4	1.00	0.75	0.86	8
5	0.99	1.00	0.99	231
6	1.00	0.99	1.00	452
accuracy			1.00	5360
macro avg	1.00	0.96	0.98	5360
weighted avg	1.00	1.00	1.00	5360

Accuracy of LightGBM: 0.9973880597014926  
 Precision of LightGBM: 0.9973967119612819  
 Recall of LightGBM: 0.9973880597014926  
 Average F1 of LightGBM: 0.9973612766651317  
 F1 of LightGBM for each type of attack: [0.99836021 0.99481865 1.099754702 0.85714286 0.99354839 0.99555556]



CPU times: user 18.8 s, sys: 184 ms, total: 19 s  
 Wall time: 10 s

```

%%time
# Train the XGBoost algorithm
import xgboost as xgb
xg = xgb.XGBClassifier()

X_train_x = X_train.values
X_test_x = X_test.values

xg.fit(X_train_x, y_train)

y_pred = xg.predict(X_test_x)
print(classification_report(y_test,y_pred))
print("Accuracy of XGBoost: "+ str(accuracy_score(y_test, y_pred)))
print("Precision of XGBoost: "+ str(precision_score(y_test, y_pred,
average='weighted'))))
print("Recall of XGBoost: "+ str(recall_score(y_test, y_pred,
average='weighted'))))
print("Average F1 of XGBoost: "+ str(f1_score(y_test, y_pred,
average='weighted'))))
print("F1 of XGBoost for each type of attack: "+ str(f1_score(y_test,
y_pred, average=None)))
xg_f1=f1_score(y_test, y_pred, average=None)

# Plot the confusion matrix
cm=confusion_matrix(y_test,y_pred)
f,ax=plt.subplots(figsize=(5,5))
sns.heatmap(cm,annot=True,linewidth=0.5,linecolor="red",fmt=".0f",ax=a
x)
plt.xlabel("y_pred")
plt.ylabel("y_true")
plt.show()

```

	precision	recall	f1-score	support
0	0.97	0.99	0.98	3656
1	0.99	0.96	0.97	387
2	0.87	0.93	0.90	14
3	0.96	0.98	0.97	612
4	1.00	0.50	0.67	8
5	0.99	0.99	0.99	231
6	0.99	0.87	0.92	452
accuracy			0.98	5360
macro avg	0.97	0.89	0.91	5360
weighted avg	0.98	0.98	0.98	5360


```

Accuracy of XGBoost: 0.9755597014925373
Precision of XGBoost: 0.9758482606600912
Recall of XGBoost: 0.9755597014925373
Average F1 of XGBoost: 0.9751364180942107

```



F1 of XGBoost for each type of attack: [0.98264171 0.97375328  
0.89655172 0.96920583 0.66666667 0.98920086  
0.92434988]



CPU times: user 38.7 s, sys: 175 ms, total: 38.9 s  
Wall time: 39 s

```
%%time
# Train the CatBoost algorithm
import catboost as cbt
cb = cbt.CatBoostClassifier(verbose=0,boosting_type='Plain')
#cb = cbt.CatBoostClassifier()

cb.fit(X_train, y_train)
y_pred = cb.predict(X_test)
print(classification_report(y_test,y_pred))
print("Accuracy of CatBoost: "+ str(accuracy_score(y_test, y_pred)))
print("Precision of CatBoost: "+ str(precision_score(y_test, y_pred,
average='weighted'))))
print("Recall of CatBoost: "+ str(recall_score(y_test, y_pred,
average='weighted'))))
print("Average F1 of CatBoost: "+ str(f1_score(y_test, y_pred,
average='weighted'))))
print("F1 of CatBoost for each type of attack: "+ str(f1_score(y_test,
y_pred, average=None)))
```

```
cb_f1=f1_score(y_test, y_pred, average=None)
```

```
# Plot the confusion matrix
```

```
cm=confusion_matrix(y_test,y_pred)
```

```
f,ax=plt.subplots(figsize=(5,5))
```

```
sns.heatmap(cm,annot=True,linewidth=0.5,linecolor="red",fmt=".0f",ax=a  
x)
```

```
plt.xlabel("y_pred")
```

```
plt.ylabel("y_true")
```

```
plt.show()
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	3656
1	0.99	0.99	0.99	387
2	1.00	1.00	1.00	14
3	1.00	1.00	1.00	612
4	1.00	0.75	0.86	8
5	0.99	1.00	0.99	231
6	1.00	0.99	0.99	452
accuracy			1.00	5360
macro avg	1.00	0.96	0.98	5360
weighted avg	1.00	1.00	1.00	5360

```
Accuracy of CatBoost: 0.9966417910447761
```

```
Precision of CatBoost: 0.9966457193764396
```


```
Recall of CatBoost: 0.9966417910447761
```

```
Average F1 of CatBoost: 0.99661603515491
```

```
F1 of CatBoost for each type of attack: [0.99781241 0.99353169 1.
```

```
0.99673203 0.85714286 0.99137931
```

```
0.9944629 ]
```



CPU times: user 3min 50s, sys: 1.59 s, total: 3min 51s  
 Wall time: 2min 3s

```
%%time
from sklearn.ensemble import
RandomForestClassifier,ExtraTreesClassifier
from sklearn.metrics import
classification_report,confusion_matrix,accuracy_score,precision_recall_
fscore_support
from sklearn.metrics import f1_score,roc_auc_score
rf = RandomForestClassifier(random_state = 0)
rf.fit(X_train,y_train)
rf_score=rf.score(X_test,y_test)
y_predict=rf.predict(X_test)
y_true=y_test
print('Accuracy of RF: ' + str(rf_score))
precision,recall,fscore,none= precision_recall_fscore_support(y_true,
y_predict, average='weighted')
print('Precision of RF: '+(str(precision)))
print('Recall of RF: '+(str(recall)))
print('F1-score of RF: '+(str(fscore)))
print(classification_report(y_true,y_predict))
cm=confusion_matrix(y_true,y_predict)
f,ax=plt.subplots(figsize=(5,5))
```


```

sns.heatmap(cm,annot=True,linewidth=0.5,linecolor="red",fmt=".0f",ax=a
x)
plt.xlabel("y_pred")
plt.ylabel("y_true")
plt.show()

```

Accuracy of RF: 0.9899253731343284  
 Precision of RF: 0.9899339378801721  
 Recall of RF: 0.9899253731343284  
 F1-score of RF: 0.9898210890672914

		precision	recall	f1-score	support
	0	0.99	1.00	0.99	3656
	1	0.99	0.98	0.99	387
	2	1.00	1.00	1.00	14
	3	0.99	0.99	0.99	612
	4	1.00	0.62	0.77	8
	5	0.99	1.00	0.99	231
	6	0.99	0.95	0.97	452
accuracy				0.99	5360
macro avg		0.99	0.94	0.96	5360
weighted avg		0.99	0.99	0.99	5360




CPU times: user 4.59 s, sys: 139 ms, total: 4.72 s  
Wall time: 4.61 s  
Compiler : 125 ms

```
%%time
from sklearn.tree import DecisionTreeClassifier
dt = DecisionTreeClassifier(random_state = 0)
dt.fit(X_train,y_train)
dt_score=dt.score(X_test,y_test)
y_predict=dt.predict(X_test)
y_true=y_test
print('Accuracy of DT: ' + str(dt_score))
precision,recall,fscore,none= precision_recall_fscore_support(y_true,
y_predict, average='weighted')
print('Precision of DT: '+(str(precision)))
print('Recall of DT: '+(str(recall)))
print('F1-score of DT: '+(str(fscore)))
print(classification_report(y_true,y_predict))
cm=confusion_matrix(y_true,y_predict)
f,ax=plt.subplots(figsize=(5,5))
sns.heatmap(cm,annot=True,linewidth=0.5,linecolor="red",fmt=".0f",ax=ax)
plt.xlabel("y_pred")
plt.ylabel("y_true")
plt.show()
```

Accuracy of DT: 0.9929104477611941  
Precision of DT: 0.9928743455303658  
Recall of DT: 0.9929104477611941  
F1-score of DT: 0.9928630672159569

	precision	recall	f1-score	support
0	1.00	0.99	1.00	3656
1	0.99	0.99	0.99	387
2	0.93	0.93	0.93	14
3	0.99	1.00	0.99	612
4	0.83	0.62	0.71	8
5	0.99	1.00	0.99	231
6	0.98	0.99	0.98	452
accuracy			0.99	5360
macro avg	0.96	0.93	0.94	5360
weighted avg	0.99	0.99	0.99	5360




CPU times: user 1.33 s, sys: 68 ms, total: 1.4 s  
Wall time: 1.44 s

```
%%time
et = ExtraTreesClassifier(random_state = 0)
et.fit(X_train,y_train)
et_score=et.score(X_test,y_test)
y_predict=et.predict(X_test)
y_true=y_test
print('Accuracy of ET: ' + str(et_score))
precision,recall,fscore,none= precision_recall_fscore_support(y_true,
y_predict, average='weighted')
print('Precision of ET: '+(str(precision)))
print('Recall of ET: '+(str(recall)))
print('F1-score of ET: '+(str(fscore)))
print(classification_report(y_true,y_predict))
cm=confusion_matrix(y_true,y_predict)
f,ax=plt.subplots(figsize=(5,5))
sns.heatmap(cm,annot=True,linewidth=0.5,linecolor="red",fmt=".0f",ax=ax)
plt.xlabel("y_pred")
plt.ylabel("y_true")
plt.show()
```

Accuracy of ET: 0.9873134328358208  
Precision of ET: 0.9873029820354025  
Recall of ET: 0.9873134328358208  
F1-score of ET: 0.9872357717235625

	precision	recall	f1-score	support
0	0.99	0.99	0.99	3656
1	0.96	0.97	0.97	387
2	0.81	0.93	0.87	14
3	0.98	0.99	0.99	612
4	0.80	0.50	0.62	8
5	0.99	1.00	0.99	231
6	0.98	0.97	0.98	452
accuracy			0.99	5360
macro avg	0.93	0.91	0.91	5360
weighted avg	0.99	0.99	0.99	5360



CPU times: user 3.68 s, sys: 135 ms, total: 3.81 s  
Wall time: 3.81 s